

ICP 6

The screenshot shows a Jupyter Notebook interface titled "ICP 6". The code cell contains Python code for creating an autoencoder using TensorFlow Keras. The code imports numpy, tensorflow.keras.layers, tensorflow.keras.models, and tensorflow.keras.callbacks. It defines sample data (x_train and x_test) and sets dimensions for input and encoded representations. It then defines the encoder (Dense layer with relu activation), decoder (Dense layer with sigmoid activation), and autoencoder model (Model). Finally, it compiles the autoencoder model.

```
import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping

# Sample data: let's use a simple dataset of 2D points for this example
x_train = np.random.random((800, 32)) # Training dataset with 800 samples and 32 features
x_test = np.random.random((200, 32)) # Test dataset with 200 samples and 32 features

# Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

# Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the autoencoder model
```

The screenshot shows the same Jupyter Notebook interface with the addition of training output. The code cell now includes the compilation of the autoencoder, definition of an EarlyStopping callback, and fitting of the autoencoder with x_train and x_test. The output shows the training progress with epochs 1 and 2, including loss values and validation loss values.

```
# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Define EarlyStopping callback
early_stopping = EarlyStopping(monitor='val_loss',
                               patience=5, # Number of epochs with no improvement
                               restore_best_weights=True) # Restores model to best weights with the lowest validation loss

# Train the autoencoder with EarlyStopping, using x_train and x_test
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                 epochs=100, # Set a high number of epochs
                 batch_size=256,
                 shuffle=True,
                 validation_data=(x_test, x_test), # Validation on the test set
                 callbacks=[early_stopping]) # Add the early stopping callback

# Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

# Let's look at the encoded representations
encoded_data = encoder.predict(x_train)
print(encoded_data.shape)
```

Epoch 1/100
4/4 2s 82ms/step - loss: 0.7074 - val_loss: 0.7035
Epoch 2/100
4/4 0s 16ms/step - loss: 0.7040 - val_loss: 0.7010

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

RAM Disk Gemini

+ Code + Text

2ds Insert code path: /tmp/100+M B

4/4 0s 17ms/step - loss: 0.6434 - val_loss: 0.6467

Epoch 100/100

4/4 0s 17ms/step - loss: 0.6449 - val_loss: 0.6478

{x} 25/25 0s 2ms/step

(800, 16)

[2] import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import TerminateOnNaN

Sample data: let's use a simple dataset of 2D points for this example
x_train = np.random.randn(800, 32) # Training dataset with 800 samples and 32 features
x_test = np.random.randn((200, 32)) # Test dataset with 200 samples and 32 features

Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

Define the input layer
input_layer = Input(shape=(input_dim,))

Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

RAM Disk Gemini

+ Code + Text

10s

encoded = Dense(encoding_dim, activation='relu')(input_layer)

Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

Define the TerminateOnNaN callback
terminate_on_nan = TerminateOnNaN()

Train the autoencoder with TerminateOnNaN, using x_train and x_test
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
epochs=100, # Set the number of epochs
batch_size=256,
shuffle=True,
validation_data=(x_test, x_test), # Validation on the test set
callbacks=[terminate_on_nan]) # Add the TerminateOnNaN callback

Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

Let's look at the encoded representations
encoded_data = encoder.predict(x_train)
print(encoded_data.shape)

✓ Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini A

RAM Disk

+ Code + Text

10s

Epoch 1/100
4/4 1s 55ms/step - loss: 0.7116 - val_loss: 0.7091
Epoch 2/100
4/4 0s 11ms/step - loss: 0.7077 - val_loss: 0.7055
Epoch 3/100
4/4 0s 11ms/step - loss: 0.7041 - val_loss: 0.7027
Epoch 4/100
4/4 0s 11ms/step - loss: 0.7020 - val_loss: 0.7004
Epoch 5/100
4/4 0s 13ms/step - loss: 0.7002 - val_loss: 0.6986
Epoch 6/100
4/4 0s 11ms/step - loss: 0.6980 - val_loss: 0.6972
Epoch 7/100
4/4 0s 11ms/step - loss: 0.6971 - val_loss: 0.6961
Epoch 8/100
4/4 0s 11ms/step - loss: 0.6959 - val_loss: 0.6951
Epoch 9/100
4/4 0s 15ms/step - loss: 0.6950 - val_loss: 0.6943
Epoch 10/100
4/4 0s 12ms/step - loss: 0.6942 - val_loss: 0.6936
Epoch 11/100
4/4 0s 11ms/step - loss: 0.6936 - val_loss: 0.6931
Epoch 12/100
4/4 0s 11ms/step - loss: 0.6930 - val_loss: 0.6926
Epoch 13/100
4/4 0s 12ms/step - loss: 0.6924 - val_loss: 0.6921
Epoch 14/100
4/4 0s 11ms/step - loss: 0.6919 - val_loss: 0.6917
Epoch 15/100
4/4 0s 11ms/step - loss: 0.6914 - val_loss: 0.6913

Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini A

RAM Disk

+ Code + Text

10s

[2] 4/4 0s 12ms/step - loss: 0.6363 - val_loss: 0.6390
Epoch 100/100
4/4 0s 11ms/step - loss: 0.6358 - val_loss: 0.6385
25/25 0s 1ms/step
(800, 16)

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import ModelCheckpoint

Sample data: let's use a simple dataset of 2D points for this example
x_train = np.random.random((800, 32)) # Training dataset with 800 samples and 32 features
x_test = np.random.random((200, 32)) # Test dataset with 200 samples and 32 features

Define the dimensions of the input and the encoded representation
input_dim = x_train.shape[1]
encoding_dim = 16 # Compress to 16 features

Define the input layer
input_layer = Input(shape=(input_dim,))

Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
✓ 78
  decoded = Dense(input_dim, activation='sigmoid')(encoded)

  # Combine the encoder and decoder into an autoencoder model
  autoencoder = Model(input_layer, decoded)

{x}
  # Compile the autoencoder model
  autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

  # Define the ModelCheckpoint callback
  checkpoint = ModelCheckpoint(filepath='autoencoder_best.keras', # File path to save the model
                               monitor='val_loss', # Metric to monitor
                               save_best_only=True, # Save only the best model based on the monitored metric
                               mode='min', # Minimize the monitored metric (e.g., validation loss)
                               save_weights_only=False, # Save the entire model (set True to save only weights)
                               verbose=1) # Print a message when saving the model

  # Train the autoencoder with the ModelCheckpoint callback
  autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                  epochs=50, # Number of epochs
                  batch_size=256,
                  shuffle=True,
                  validation_data=(x_test, x_test), # Validation data
                  callbacks=[checkpoint]) # Add the ModelCheckpoint callback

  # Once trained, we can use the encoder part separately
  encoder = Model(input_layer, encoded)

  # Let's look at the encoded representations
```

✓ Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
✓ 78
  # Once trained, we can use the encoder part separately
  encoder = Model(input_layer, encoded)

{x}
  # Let's look at the encoded representations
  encoded_data = encoder.predict(x_train)
  print(encoded_data.shape)

  Epoch 1/50
  1/4 ━━━━━━ 2s 855ms/step - loss: 0.7062
  Epoch 1: val_loss improved from inf to 0.70645, saving model to autoencoder_best.keras
  4/4 ━━━━━━ 1s 72ms/step - loss: 0.7066 - val_loss: 0.7064
  Epoch 2/50
  1/4 ━━━━ 0s 26ms/step - loss: 0.7037
  Epoch 2: val_loss improved from 0.70645 to 0.70376, saving model to autoencoder_best.keras
  4/4 ━━━━ 0s 17ms/step - loss: 0.7039 - val_loss: 0.7038
  Epoch 3/50
  1/4 ━━━━ 0s 89ms/step - loss: 0.7024
  Epoch 3: val_loss improved from 0.70376 to 0.70155, saving model to autoencoder_best.keras
  4/4 ━━━━ 0s 27ms/step - loss: 0.7019 - val_loss: 0.7015
  Epoch 4/50
  1/4 ━━━━ 0s 58ms/step - loss: 0.7004
  Epoch 4: val_loss improved from 0.70155 to 0.69976, saving model to autoencoder_best.keras
  4/4 ━━━━ 0s 18ms/step - loss: 0.6999 - val_loss: 0.6998
  Epoch 5/50
  1/4 ━━━━ 0s 82ms/step - loss: 0.6987
  Epoch 5: val_loss improved from 0.69976 to 0.69832, saving model to autoencoder_best.keras
  4/4 ━━━━ 0s 19ms/step - loss: 0.6984 - val_loss: 0.6983
  Epoch 6/50
```

✓ Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini

+ Code + Text

7s [4] Epoch 49: val_loss improved from 0.67061 to 0.66992, saving model to autoencoder_best.keras
4/4 0s 17ms/step - loss: 0.6688 - val_loss: 0.6699

[x] Epoch 50/50
1/4 0s 87ms/step - loss: 0.6686
Epoch 50: val_loss improved from 0.66992 to 0.66920, saving model to autoencoder_best.keras
4/4 0s 18ms/step - loss: 0.6681 - val_loss: 0.6692
25/25 0s 2ms/step
(800, 16)

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import ReduceLROnPlateau

Sample data: let's use a simple dataset of 2D points
data = np.random.random((1000, 32)) # 1000 samples of 32 features

Define the dimensions of the input and the encoded representation
input_dim = data.shape[1]
encoding_dim = 16 # Compress to 16 features

Define the input layer
input_layer = Input(shape=(input_dim,))

Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

Define the decoder

✓ Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini

+ Code + Text

5s [x] # Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

Define the ReduceLROnPlateau callback
reduce_lr = ReduceLROnPlateau(monitor='val_loss', # Metric to monitor
 factor=0.5, # Factor by which the learning rate will be reduced (new_lr = lr * factor)
 patience=3, # Number of epochs with no improvement after which learning rate will be reduced
 min_lr=1e-6, # Lower bound for the learning rate
 verbose=1) # Print message when the learning rate is reduced

Assuming x_train and x_test are your training and validation datasets
Train the autoencoder with the ReduceLROnPlateau callback
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
 epochs=50, # Number of epochs
 batch_size=256,
 shuffle=True,
 validation_data=(x_test, x_test), # Validation data
 callbacks=[reduce_lr]) # Add the ReduceLROnPlateau callback

Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

✓ Connected to Python 3 Google Compute Engine backend

ICP 6

File Edit View Insert Runtime Tools Help All changes saved

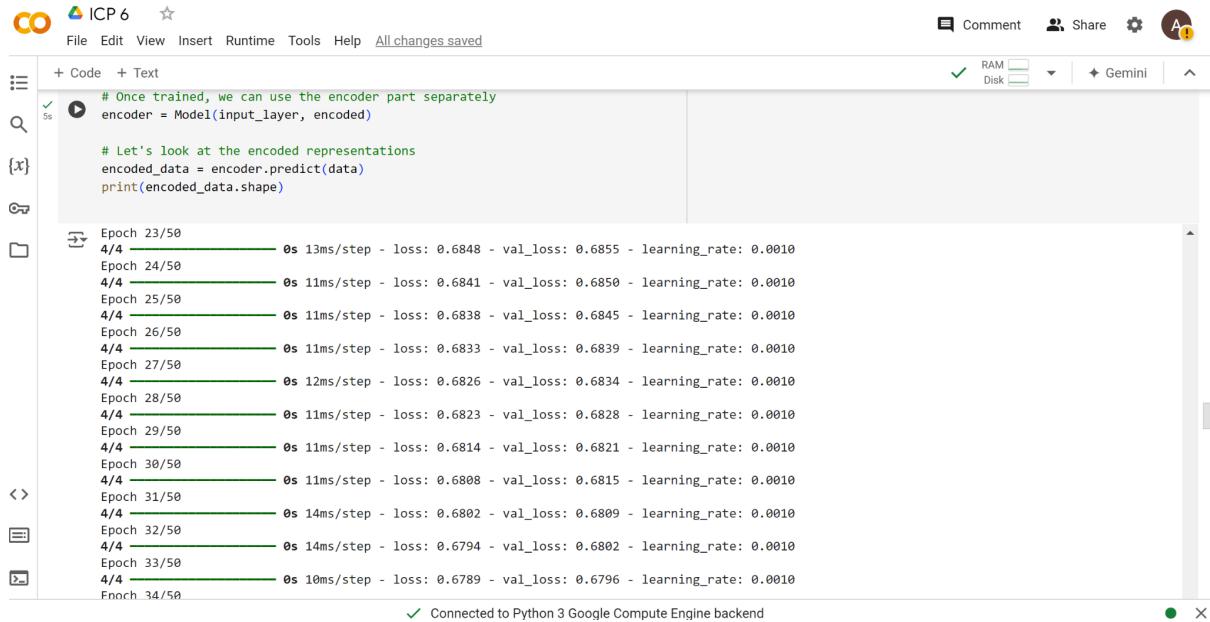
+ Code + Text

```
# Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

# Let's look at the encoded representations
encoded_data = encoder.predict(data)
print(encoded_data.shape)
```

RAM Disk Gemini

Connected to Python 3 Google Compute Engine backend



ICP 6

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[5] 4/4 ━━━━━━ 0s 11ms/step - loss: 0.6684 - val_loss: 0.6692 - learning_rate: 0.0010
Epoch 50/50
4/4 ━━━━━━ 0s 11ms/step - loss: 0.6674 - val_loss: 0.6686 - learning_rate: 0.0010
32/32 ━━━━━━ 0s 2ms/step
(1000, 16)

import numpy as np
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, TerminateOnNaN, ReduceLROnPlateau

# Sample data: let's use a simple dataset of 2D points
data = np.random.random((1000, 32)) # 1000 samples of 32 features

# Define the dimensions of the input and the encoded representation
input_dim = data.shape[1]
encoding_dim = 16 # Compress to 16 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

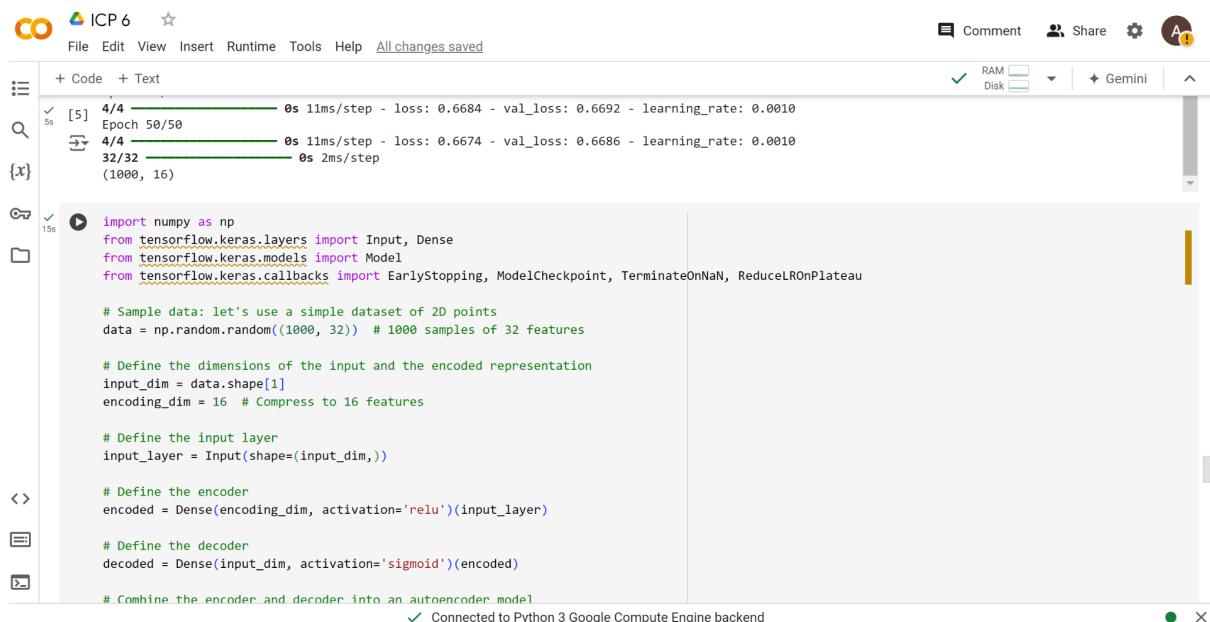
# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

# Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

# Combine the encoder and decoder into an autoencoder model
```

RAM Disk Gemini

Connected to Python 3 Google Compute Engine backend



ICP 6

```

15s # Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

# Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

# Combine the encoder and decoder into an autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the autoencoder model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Define the callbacks
# EarlyStopping to stop training if validation loss doesn't improve
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# ModelCheckpoint to save the best model based on validation loss
checkpoint = ModelCheckpoint(filepath='autoencoder_best.keras', monitor='val_loss', save_best_only=True, verbose=1)

# TerminateOnNaN to stop training if the loss becomes NaN
terminate_on_nan = TerminateOnNaN()

# ReduceLROnPlateau to reduce the learning rate when validation loss plateaus
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6, verbose=1)

# Assuming x_train and x_test are your training and validation datasets
# Training the autoencoder with all the callbacks
autoencoder.fit(x_train, x_train, # For autoencoders input and output are the same
                epochs=100, batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test), # Validation data
                callbacks=[reduce_lr, early_stopping, checkpoint, terminate_on_nan]) # Using multiple callbacks

# Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

# Let's look at the encoded representations
encoded_data = encoder.predict(data)
print(encoded_data.shape)

✓ Connected to Python 3 Google Compute Engine backend

```

ICP 6

```

15s # Assuming x_train and x_test are your training and validation datasets
# Training the autoencoder with all the callbacks
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                epochs=100, # Set a higher number of epochs, since callbacks will handle stopping early
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test), # Validation data
                callbacks=[reduce_lr, early_stopping, checkpoint, terminate_on_nan]) # Using multiple callbacks

# Once trained, we can use the encoder part separately
encoder = Model(input_layer, encoded)

# Let's look at the encoded representations
encoded_data = encoder.predict(data)
print(encoded_data.shape)

Epoch 1/100
1/4 2s 798ms/step - loss: 0.7144
Epoch 1: val_loss improved from inf to 0.71033, saving model to autoencoder_best.keras
4/4 1s 60ms/step - loss: 0.7139 - val_loss: 0.7103 - learning_rate: 0.0010
Epoch 2/100
1/4 0s 28ms/step - loss: 0.7093
Epoch 2: val_loss improved from 0.71033 to 0.70587, saving model to autoencoder_best.keras
4/4 0s 18ms/step - loss: 0.7086 - val_loss: 0.7059 - learning_rate: 0.0010
Epoch 3/100
1/4 0s 85ms/step - loss: 0.7042
Epoch 3: val_loss improved from 0.70587 to 0.70254, saving model to autoencoder_best.keras
4/4 0s 17ms/step - loss: 0.7041 - val_loss: 0.7025 - learning rate: 0.0010
✓ Connected to Python 3 Google Compute Engine backend

```

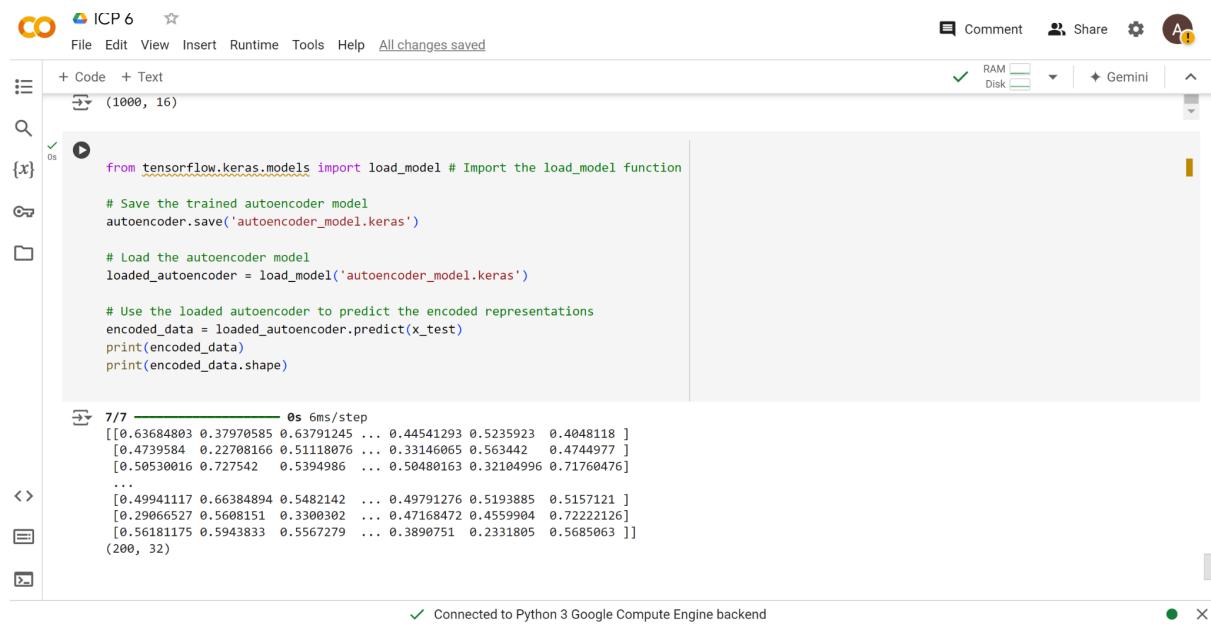
ICP 6

```

File Edit View Insert Runtime Tools Help All changes saved
15s [7] Epoch 99: val_loss improved from 0.64151 to 0.64111, saving model to autoencoder_best.keras
4/4 0s 18ms/step - loss: 0.6413 - val_loss: 0.6411 - learning_rate: 0.0010
    Epoch 100/100
1/4 0s 29ms/step - loss: 0.6399
Epoch 100: val_loss improved from 0.64111 to 0.64071, saving model to autoencoder_best.keras
4/4 0s 17ms/step - loss: 0.6401 - val_loss: 0.6407 - learning_rate: 0.0010
32/32 0s 2ms/step
(1000, 16)

```

Loading the model and predicting



The screenshot shows a Google Colab notebook titled "ICP 6". The code cell contains Python code for loading a trained autoencoder and using it to predict encoded representations from test data. The output of the code shows the prediction results and their shape.

```
from tensorflow.keras.models import load_model # Import the load_model function

# Save the trained autoencoder model
autoencoder.save('autoencoder_model.keras')

# Load the autoencoder model
loaded_autoencoder = load_model('autoencoder_model.keras')

# Use the loaded autoencoder to predict the encoded representations
encoded_data = loaded_autoencoder.predict(x_test)
print(encoded_data)
print(encoded_data.shape)

7/7 0s 6ms/step
[[0.63684803 0.37970585 0.63791245 ... 0.44541293 0.5235923 0.4048118 ]
 [0.4739584 0.22708166 0.51118076 ... 0.33146065 0.563442 0.4744977 ]
 [0.50530016 0.727542 0.5394986 ... 0.50480163 0.32104996 0.71760476]
 ...
 [0.49941117 0.66384894 0.5482142 ... 0.49791276 0.5193885 0.5157121 ]
 [0.29066527 0.5608151 0.3300302 ... 0.47168472 0.4559984 0.72222126]
 [0.56181175 0.5943833 0.5567279 ... 0.3890751 0.2331805 0.5685063 ]]
(200, 32)
```

Connected to Python 3 Google Compute Engine backend

Youtube video:<https://youtu.be/DYkWaphYpko>

GitHub link :<https://github.com/w8162583/bda.git>